

Lab#3

October 27, 2022

0.1 Lab 3: Wide data and linear models

0.1.1 Alice Wu. ID: qichaow

```
[131]: ## Import libraries
import numpy as np
import pandas as pd
import seaborn as sns
import os
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn import manifold
from matplotlib import pyplot

from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression, LinearRegression, Lasso, Ridge, ElasticNet
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import precision_score, recall_score, f1_score
from sklearn.model_selection import train_test_split
```

```
[132]: testX=pd.read_csv('/Users/aliceqichaowu/Desktop/38615/test_X.csv')
trainX=pd.read_csv('/Users/aliceqichaowu/Desktop/38615/train_X.csv')
trainY=pd.read_csv('/Users/aliceqichaowu/Desktop/38615/train_y.csv')
```

```
[133]: print(testX.info())
print(trainX.info())
print(trainY.info())
testX.head(5)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 110 entries, 0 to 109
Columns: 17971 entries, Unnamed: 0 to ENSG00000283697
dtypes: float64(17970), object(1)
memory usage: 15.1+ MB
None
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 444 entries, 0 to 443
Columns: 17971 entries, Unnamed: 0 to ENSG000000283697
dtypes: float64(17970), object(1)
memory usage: 60.9+ MB
None
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 444 entries, 0 to 443
Data columns (total 2 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            444 non-null    object
1   xml_neoplasm_histologic_grade         444 non-null    int64
dtypes: int64(1), object(1)
memory usage: 7.1+ KB
None

```

```

[133]:
           Unnamed: 0  ENSG000000000003  ENSG000000000005  \
0  1576BB97-F8EA-48FF-9494-EBC7A0C97312    7.503849    1.685440
1  63F4281A-4D90-4589-879E-549097FB3459    6.112919   -5.425877
2  74CF740D-201D-4070-99B9-F007E7C4D472    6.183846    1.217355
3  1329E1C1-743E-4541-9081-165572D4BF47    6.325535   -0.197432
4  88BBA1BF-44E0-416E-A028-04F410A3FE5D    7.162383   -5.539710

           ENSG000000000419  ENSG000000000457  ENSG000000000938  ENSG000000000971  \
0           8.090089           5.546588           3.412025           5.871539
1           5.604743           4.248246           3.374033           3.370075
2           6.093903           4.403216           5.722867           6.123209
3           6.722632           4.509093           4.941256           3.700171
4           6.186110           5.270282           5.412103           3.781568

           ENSG000000001036  ENSG000000001084  ENSG000000001167  ...  ENSG000000282651  \
0           8.395029           5.894996           6.313896  ...           3.626396
1           6.804015           4.646287           6.236134  ...          -5.425877
2           7.680258           4.983386           5.626569  ...          11.317818
3           7.492606           5.031053           5.859242  ...           3.536141
4           7.285779           4.847552           6.337205  ...           5.112772

           ENSG000000282815  ENSG000000282939  ENSG000000283063  ENSG000000283439  \
0          -2.640463          -7.278265           0.721287          -7.278265
1          -2.143410           1.956699           1.377768          -5.425877
2          -5.259782           5.093316           3.800607          -5.259782
3          -1.362703           3.480703           0.776383          -0.999629
4          -5.539710           4.369268           2.927849           0.791397

           ENSG000000283463  ENSG000000283526  ENSG000000283586  ENSG000000283632  \
0           4.672536           5.153675          -7.278265           4.843800

```

1	4.560829	3.735085	-5.425877	4.105789
2	4.966710	2.829487	-5.259782	5.118704
3	4.794226	-5.399485	-5.399485	4.660878
4	4.542333	4.628775	-5.539710	4.812787

	ENSG00000283697
0	4.302864
1	4.803357
2	5.024979
3	4.684343
4	4.524834

[5 rows x 17971 columns]

```
[231]: # Check missing values
print('There are %i nan in the dataframe' % trainY.isna().sum().sum())
```

There are 0 nan in the dataframe

```
[134]: # Set up X_test, X_train, and y
y = trainY.iloc[:,1:]
X_train = trainX.iloc[:,1:]
X_test= testX.iloc[:,1:]
```

```
[135]: # Normalization
scaler = StandardScaler() # normalization: zero mean, unit variance
scaler.fit(X_train) # scaling factor determined from the training set

X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test) # apply the same scaling to the test set
```

0.1.2 1. Try different linear models: linear regression, logistic regression, Ridge regression, LASSO

```
[51]: ## Linear regression:
# accruacy: 0.91666 (private), 0.84375 (public)
mdl_lar= LinearRegression()
mdl_lar.fit(X_train, y)
y_pred_LaR_temp = np.squeeze(mdl_lar.predict(X_test))
y_pred_LaR=[]
for i in y_pred_LaR_temp:
    if i>0.5:
        y_pred_LaR.append(1)
    else:
        y_pred_LaR.append(0)
# [ x for x in y_pred_LaR if x != 0 ]
```

```
[75]: ## Logistic regression
# accruacy: 0.88888 (private), 0.83870 (public)
mdl_lr = LogisticRegression('none')
mdl_lr.fit(X_train, y)
y_pred_lr_temp = np.squeeze(mdl_lr.predict(X_test))
y_pred_LR=[]
for i in y_pred_lr_temp:
    if i>0.5:
        y_pred_LR.append(1)
    else:
        y_pred_LR.append(0)
```

/Users/aliceqichaowu/opt/anaconda3/lib/python3.8/site-packages/sklearn/utils/validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

```
[65]: ## Ridge regression
# accruacy: 0.91666 (private), 0.84375 (public)
mdl_rr = Ridge(alpha=0.1)
mdl_rr.fit(X_train, y)
y_pred_rr_temp = np.squeeze(mdl_rr.predict(X_test))
y_pred_RR=[]
for i in y_pred_rr_temp:
    if i>0.5:
        y_pred_RR.append(1)
    else:
        y_pred_RR.append(0)
```

```
[66]: ## Lasso regression:
# accruacy: 0.87671 (private), 0.86567 (public)
mdl_las = Lasso(alpha=0.1)
mdl_las.fit(X_train, y)
y_pred_las_temp = np.squeeze(mdl_las.predict(X_test))
y_pred_Las=[]
for i in y_pred_las_temp:
    if i>0.5:
        y_pred_Las.append(1)
    else:
        y_pred_Las.append(0)
```

```
[73]: ## Elastic Net
# accruacy: 0.87323 (private), 0.84848 (public)
mdl_en = ElasticNet(alpha=0.1)
mdl_en.fit(X_train, y)
y_pred_en_temp = np.squeeze(mdl_en.predict(X_test))
y_pred_EN=[]
```

```

for i in y_pred_en_temp:
    if i>0.5:
        y_pred_EN.append(1)
    else:
        y_pred_EN.append(0)

```

0.1.3 2. Test regularization parameters:

- Try L1 (lasso) and L2 (ridge) on logistic regression; Try lasso and ridge regression models with with different alpha values
- After trying different models, I found out the logistic model with penalty='elasticnet', solver='saga', l1_ratio=0.5 has the highest accuracy 0.90140 (private score) and 0.88888 (public score).

```

[165]: ## Logistic regression
# LR1 (logistic, penalty='l1', solver='liblinear'): 0.85714 (private), 0.81967 (public)
# LR2 (logistic, penalty='elasticnet', solver='saga', l1_ratio=0.5): 0.90140, 0.88888
# LR2, alpha=100: 0.79120, 0.76404

mdl_lr1 = LogisticRegression(penalty='l1', solver='liblinear')
mdl_lr2 = LogisticRegression(penalty='elasticnet', solver='saga', l1_ratio=0.5)
mdl_en_a1 = ElasticNet(alpha=1)
mdl_en_a100 = ElasticNet(alpha=100)

mdl_lr1.fit(X_train, y)
mdl_lr2.fit(X_train, y)
mdl_en_a1.fit(X_train, y)
mdl_en_a100.fit(X_train, y)

y_pred_lr_temp1 = np.squeeze(mdl_lr1.predict(X_test))
y_pred_lr_temp2 = np.squeeze(mdl_lr2.predict(X_test))
y_pred_en_a1_temp = np.squeeze(mdl_en_a1.predict(X_test))
y_pred_en_a100_temp = np.squeeze(mdl_en_a100.predict(X_test))

y_pred_LR1 = []
y_pred_LR2 = []
y_pred_en_a1 = []
y_pred_en_a100 = []

for i in y_pred_en_a100_temp:
    if i>0.5:
        y_pred_en_a100.append(1)
    else:
        y_pred_en_a100.append(0)

```

```

/Users/aliceqichaowu/opt/anaconda3/lib/python3.8/site-
packages/sklearn/utils/validation.py:63: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    return f(*args, **kwargs)
/Users/aliceqichaowu/opt/anaconda3/lib/python3.8/site-
packages/sklearn/utils/validation.py:63: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    return f(*args, **kwargs)
/Users/aliceqichaowu/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_sag.py:328: ConvergenceWarning: The max_iter was
reached which means the coef_ did not converge
    warnings.warn("The max_iter was reached which means ")

```

0.1.4 3. Random shuffle

- Model accuracy: 0.67605 (private), 0.65671 (public)

We can see that the random model will have a random chance level

```

[130]: from sklearn.utils import shuffle
X_ran, y_ran = shuffle(X_train, y)

reg = LinearRegression()
reg.fit(X_ran, y_ran)

y_pred_ran_temp = reg.predict(X_test)
y_pred_ran = []
for i in y_pred_ran_temp:
    if i > 0.5:
        y_pred_ran.append(1)
    else:
        y_pred_ran.append(0)
# y_pred_ran

```

```

[152]: ## export the y_pred data file
y_pred_col = pd.DataFrame(data={'ID': testX.iloc[:,
    ↪ 0], 'xml_neoplasm_histologic_grade': y_pred_en_a100})
os.makedirs('/Users/aliceqichaowu/Desktop/38615/', exist_ok=True)
y_pred_col.to_csv('/Users/aliceqichaowu/Desktop/38615/pred_en_a100.
    ↪ csv', index=False)

```

0.1.5 4. Feature importance: find out the 10 most important features

```

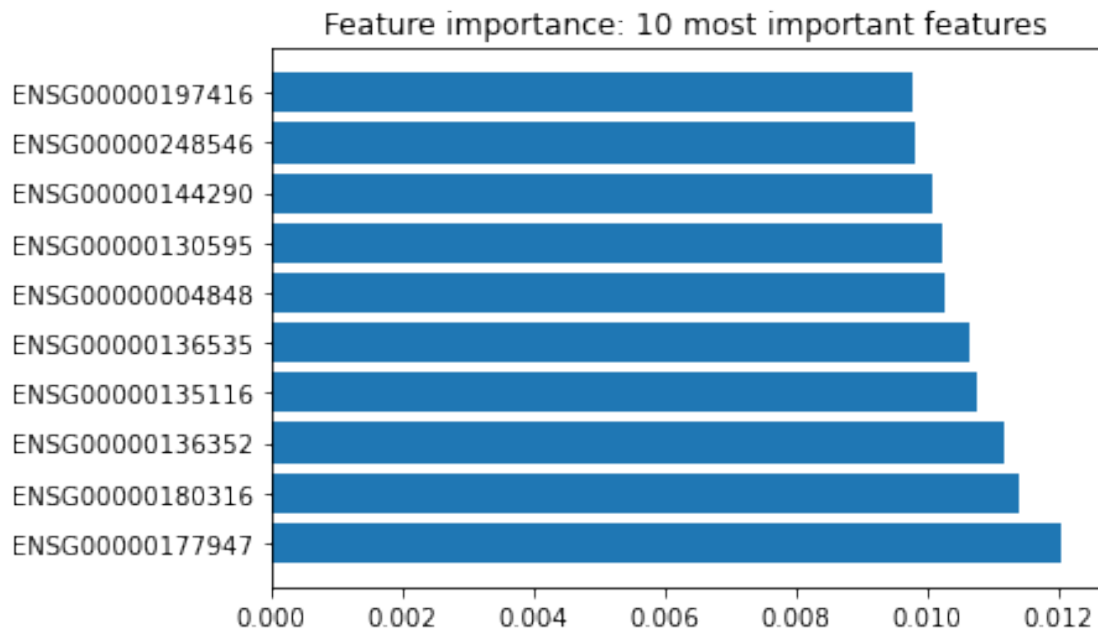
[224]: importance = np.flip(np.argsort mdl_lr2.coef_[0]))
importance10 = importance[0:10]
all_feat = testX.columns[1:]

```

```
all_arr=np.array(all_feat)
print(all_arr[importance10])
plt.barh(all_arr[importance10], np.flip(np.sort mdl_lr2.coef_[0]))[0:10])
plt.title('Feature importance: 10 most important features')
```

```
['ENSG000000177947' 'ENSG000000180316' 'ENSG000000136352' 'ENSG000000135116'
 'ENSG000000136535' 'ENSG000000004848' 'ENSG000000130595' 'ENSG000000144290'
 'ENSG000000248546' 'ENSG000000197416']
```

[224]: Text(0.5, 1.0, 'Feature importance: 10 most important features')




0.1.6 5. Visualize the dataset: implement tSNE to visualize the clusters

- Here, we can see it is hard to differentiate two classes with tSNE

```
[234]: tsne = manifold.TSNE(random_state=42,n_components=2)
X_tsne = tsne.fit_transform(X_train)
colors = {0:'b', 1:'r'}
fig1,ax = plt.subplots()
ax.scatter(X_tsne[:,0], X_tsne[:,1],c=trainY.iloc[:,-1].map(colors), s=2)
ax.set_xlabel('tSNE 1')
ax.set_ylabel('tSNE 2')
plt.title('Training data visualized by tSNE')
```

[234]: Text(0.5, 1.0, 'Training data visualized by tSNE')



[]: